

Matrix Completion Methods for Causal Panel Data Models

Susan Athey et al. (2021), JASA

Naoki Eguchi

2025.6.25 ミクロ計量経済学

Faculty of Medicine, Kyoto University

Introduction

What is "matrix completion" ?

- Matrix completion (MC) is an imputation method for the missing in the matrix.
- In PCA (Principle Component Analysis), given M principle components, we approximate a true data matrix $\mathbf{Y} \in \mathbb{R}^{I \times J}$ by $\mathbf{A} \in \mathbb{R}^{I \times M}$ and $\mathbf{B} \in \mathbb{R}^{M \times J}$.

$$\min_{\mathbf{A}, \mathbf{B}} \left\{ \sum_{i=1}^I \sum_{j=1}^J \left(Y_{ij} - \sum_{m=1}^M a_{im} b_{mj} \right)^2 \right\}$$

- If there is missing in matrix, using only observed data, we can approximate the original matrix, hence we can impute the missing by $\hat{\mathbf{A}}, \hat{\mathbf{B}}$.

$$\min_{\mathbf{A}, \mathbf{B}} \left\{ \sum_{(i,j) \in \mathcal{O}} \left(Y_{ij} - \sum_{m=1}^M a_{im} b_{mj} \right)^2 \right\}$$

- Applying this idea to panel data analysis, MC-based panel method is seen to a generalization of many model-based (factor-regression-based) panel methods.

Where to impute


- As many panel data methods, we want to know ATT : $\mathbb{E}[Y_{it}(1) - Y_{it}(0)|W_i = 1]$.
- Thus, it boils down to estimate (impute) the counterfactual $Y_{it}(0)$.
 - Horizontal : Under unconfoundedness, we can impute counterfactual PO using observed outcomes for control units.
 - Vertical : By SCM, we can also impute it using weighted average outcomes for control units with most predictive weights trained with pre-treatment datas.

$$\mathbf{Y} = \begin{pmatrix} \checkmark & \checkmark & \checkmark \\ \vdots & \vdots & \vdots \\ \checkmark & \checkmark & \checkmark \\ \checkmark & \checkmark & ? \\ \vdots & \vdots & \vdots \\ \checkmark & \checkmark & ? \end{pmatrix}.$$

$$\mathbf{Y} = \begin{pmatrix} \checkmark & \checkmark & \dots & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \dots & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \dots & \checkmark & ? & \dots & ? \end{pmatrix}$$

Xu (2024): Counterfactual estimation

- functional form: $Y_{it}(0) = f(\mathbf{X}_{it}) + h(\mathbf{U}_{it}) + \epsilon_{it}$
→ No anticipation, carryover, feedback, LDV
- strict exogeneity: $\forall i, j \in \{1, \dots, N\}, \forall s, t \in \{1, \dots, T\}, \epsilon_{it} \perp\!\!\!\perp \{D_{js}, \mathbf{X}_{js}, \mathbf{U}_{js}\}$



- low-dimensional decomposition: $h(\mathbf{U}_{it}) = \{L_{it}\}, \text{rank}(\mathbf{L}_{N \times T}) \ll \min\{N, T\}$
→ The rank (= number of factors) is **FIXED !!**
- MC panel method is a generalization of Xu's counterfactual estimation.

Set Up

Notation and Estimand

- Consider a setting with N units observed over T periods characterized by a binary treatment W_{it} and hence two POs $Y_{it}(1), Y_{it}(0)$.
- Estimand : $\mathbf{Y} = \{Y_{it}(0)^1\} = \begin{pmatrix} Y_{11}(0) & \cdots & Y_{1T}(0) \\ \vdots & \ddots & \vdots \\ Y_{N1}(0) & \cdots & Y_{NT}(0) \end{pmatrix}$ (\leftarrow Matrix!!)
- $W_{it} = \begin{cases} 1 & \text{if } (i, t) \in \mathcal{M} : \text{Missing indice} \\ 0 & \text{if } (i, t) \in \mathcal{O} : \text{Observed indice as training data} \end{cases}$

¹以降は簡単のため, $Y_{it}(0) = Y_{it}$ とし, “(0)” を省略して表記する.

Patterns of data matrix

- Ordinary case (rich data wrt. units and times)

$$\mathbf{Y}_{N \times T} = \begin{pmatrix} \checkmark & \checkmark & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \checkmark & ? & \dots & ? \\ \checkmark & \checkmark & \checkmark & ? & \dots & ? \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ \checkmark & \checkmark & \checkmark & ? & \dots & ? \end{pmatrix}$$

- Staggered adoption

$$\mathbf{Y}_{N \times T} = \begin{pmatrix} \checkmark & \checkmark & \checkmark & \checkmark & \dots & \checkmark & \text{(never adopter)} \\ \checkmark & \checkmark & \checkmark & \checkmark & \dots & ? & \text{(late adopter)} \\ \checkmark & \checkmark & ? & ? & \dots & ? & \\ \checkmark & \checkmark & ? & ? & \dots & ? & \text{(medium adopter)} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \\ \checkmark & ? & ? & ? & \dots & ? & \text{(early adopter)} \end{pmatrix}$$

Horizontal regression and unconfoundedness : thin matrix ($N \gg T$)

$$\mathbf{Y} = \begin{pmatrix} \checkmark & \checkmark & \checkmark \\ \vdots & \vdots & \vdots \\ \checkmark & \checkmark & \checkmark \\ \checkmark & \checkmark & ? \\ \vdots & \vdots & \vdots \\ \checkmark & \checkmark & ? \end{pmatrix}.$$

1. Regress the last period outcome on the lagged outcomes. (among untreated)
2. Predict the missing POs using the estimated regression.

$$\forall (i, T) \in \mathcal{M}, \hat{Y}_{iT} = \hat{\beta}_0 + \sum_{t=1}^{T-1} \hat{\beta}_t Y_{it}, \text{ where } \hat{\beta} = \arg \min_{\beta} \sum_{i:(i,T) \in \mathcal{O}} (Y_{iT} - \beta_0 - \sum_{t=1}^{T-1} \beta_t Y_{it})^2.$$

→ To make this more flexible and nonparametric, it leads to matching where we find for each treated unit i a corresponding unit j to approximate Y_{it} by Y_{jt} for $t = 1, \dots, T - 1$.

Vertical regression and synthesis control : fat matrix ($T \gg N$)

$$\mathbf{Y} = \begin{pmatrix} \checkmark & \checkmark & \dots & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \dots & \checkmark & \checkmark & \dots & \checkmark \\ \checkmark & \checkmark & \dots & \checkmark & ? & \dots & ? \end{pmatrix}$$

1. Regress the outcomes for treated unit prior to the treatment on the outcomes for the control units in the same periods.
2. Predict the missing POs using the estimated regression.

$$\forall (N, t) \in \mathcal{M}, \hat{Y}_{Nt} = \hat{\gamma}_0 + \sum_{i=1}^{N-1} \hat{\gamma}_i Y_{it}, \text{ where } \hat{\gamma} = \arg \min_{\gamma} \sum_{t:(N,t) \in \mathcal{O}} (Y_{Nt} - \gamma_0 - \sum_{i=1}^{N-1} \gamma_i Y_{it})^2.$$

→ Vertical regression is generalization of ADH(2010) in that it relaxes two restrictions :

- the coefficients $\hat{\gamma}$ are nonnegative. (Interpretability ; What is a negative weight?)
- the intercept in this regression is 0. (This is seen to be plausible in recent literatures.)

Matrix Completion

- Under no covariates, we model the $N \times T$ matrix of complete matrix \mathbf{Y} as

$$\mathbf{Y} = \mathbf{L}^* + \epsilon, \text{ where } \mathbb{E}[\epsilon|\mathbf{L}^*] = 0.$$

Assumption 1

- ϵ is independent of \mathbf{L}^* (strict exogeneity)
- The element of ϵ are σ - *sub* - *Gaussian* and independent each other.
 $\Leftrightarrow \forall t, \mathbb{E}[\exp(t\epsilon)] \leq \exp(\frac{\sigma^2 t^2}{2}).$

- The goal is to estimate the matrix \mathbf{L}^* . (low-rank assumption)

→ Note that two types² of fixed effects are included.

²これら以外にも Interactive fixed effect といったあらゆる factor を”少数まで”許容する

MC-NNM (Matrix Completion with Nuclear Norm Minimization) estimator

- MC-NNM estimator for \mathbf{L}^* is given by $\hat{\mathbf{L}} + \hat{\Gamma}\mathbf{1}_T^\top + \mathbf{1}_N\hat{\Delta}^\top$

$$(\hat{\mathbf{L}}, \hat{\Gamma}, \hat{\Delta}) = \arg \min_{\mathbf{L}, \Gamma, \Delta} \left\{ \frac{1}{|\mathcal{O}|} \|\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \mathbf{L} - \Gamma\mathbf{1}_T^\top - \mathbf{1}_N\Delta^\top)\|_F^2 + \lambda \|\mathbf{L}\|_* \right\}$$

- $\Gamma \in \mathbb{R}^N$: unit-varying (and time-fixed) effect (individual effect)
- $\Delta \in \mathbb{R}^T$: time-varying (and unit-fixed) effect (time effect)
- matrix indicator function : $\mathbf{P}_{\mathcal{O}}(\mathbf{A}) = \begin{cases} A_{it} & \text{if } (i, t) \in \mathcal{O} \\ 0 & \text{if } (i, t) \notin \mathcal{O} \end{cases}$ (NA is regarded as 0)
- Frobenius norm : $\|\mathbf{A}\|_F^2 = \sum_{i=1}^N \sum_{t=1}^T A_{it}^2$ (行列版の mean squared error を計算している)
- Regularization term $\lambda \|\mathbf{L}\|_*$ leads to **the low rank of \mathbf{L}** .
 \rightarrow minimize $\lambda \|\mathbf{L}\|_* \Leftrightarrow$ the sparsity of Singular value $\sigma_i(\mathbf{L})(> 0) \Leftrightarrow$ low rank of \mathbf{L}

- **Fact 1.** (Singular value decomposition) *Every real matrix $L \in \mathbb{R}^N \times \mathbb{R}^T$ can be decomposed using a orthogonal matrix $\mathbf{S} \in \mathbb{R}^N \times \mathbb{R}^{\min(N,T)}$, $\mathbf{R} \in \mathbb{R}^T \times \mathbb{R}^{\min(N,T)}$ by*

$$\mathbf{L} = \mathbf{S}\mathbf{\Sigma}\mathbf{R}', \text{ where } \mathbf{\Sigma} = \text{diag}(\sigma_1, \dots, \sigma_{\min(N,T)}), \mathbf{S}'\mathbf{S} = \mathbf{I}_{\min(N,T)} = \mathbf{R}'\mathbf{R}$$

- **Fact 2.** *The number of non-zero singular value = rank \mathbf{L}*

- Nuclear norm : $\|L\|_* = \sum_{i=1}^{\min(N,T)} \sigma_i(\mathbf{L})$

\rightarrow minimize $\lambda\|L\|_* \Leftrightarrow$ the sparsity of Singular value $\sigma_i(\mathbf{L})(> 0) \Leftrightarrow$ low rank of \mathbf{L}

- Since the rank of \mathbf{L} corresponds to **the number of factor**, this assumption of low rank is quite plausible.
- Although the low rank matrix CAN include two fixed effects, these "strong" factors are separately estimated for improving the quality of the practical imputations.

Algorithm for calculating $\hat{\mathbf{L}}$

- For simplicity, assume that there are no fixed effects. (only estimate \mathbf{L})
- **Fact 3.** For $\mathbf{A} = \mathbf{S}\mathbf{\Sigma}\mathbf{R}$, the minimizer is obtained analytically.

$$\mathbf{S}\tilde{\mathbf{\Sigma}}\mathbf{R}^\top = \arg \min_{\mathbf{A}} \left\{ \frac{1}{2} \|\mathbf{L} - \mathbf{A}\|_F^2 + \lambda \|\mathbf{A}\|_* \right\}, \text{ where } \tilde{\mathbf{\Sigma}} = \text{diag}(\{\max(\sigma_i(\mathbf{A}) - \lambda, 0)\}_i)$$

→ You can see elements with a small singular value (=weak factor) will be vanished.

- We perform this minimization over and over until the matrix converges.
 - Define $\text{shrink}_\lambda(\mathbf{A}) = \mathbf{S}\tilde{\mathbf{\Sigma}}\mathbf{R}^\top$ and start with the initial choice $\mathbf{L}_1(\lambda, \mathcal{O}) = \mathbf{P}_{\mathcal{O}}(\mathbf{Y})$ (The missing starts with 0.)

$$\mathbf{L}_{k+1}(\lambda, \mathcal{O}) = \text{shrink}_{\frac{\lambda|\mathcal{O}|}{2}} \{ \mathbf{P}_{\mathcal{O}}(\mathbf{Y}) + \mathbf{P}_{\mathcal{O}}^\top(\mathbf{L}_k(\lambda, \mathcal{O})) \}$$

- $\mathbf{P}_{\mathcal{O}}(\mathbf{A}) = \begin{cases} A_{it} & \text{if } (i, t) \in \mathcal{O} \\ 0 & \text{if } (i, t) \notin \mathcal{O} \end{cases}, \quad \mathbf{P}_{\mathcal{O}}^\top(\mathbf{A}) = \begin{cases} 0 & \text{if } (i, t) \in \mathcal{O} \\ A_{it} & \text{if } (i, t) \notin \mathcal{O} \end{cases}$
- The limiting matrix $\hat{\mathbf{L}}(\lambda, \mathcal{O}) = \lim_{k \rightarrow \infty} \mathbf{L}_k(\lambda, \mathcal{O})$ is MC-NNM estimator given λ .

- For the case with fixed effects, we replace $\mathbf{P}_{\mathcal{O}}(\mathbf{Y})$ with $\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \Gamma_k \mathbf{1}_T^T - \mathbf{1}_N \Delta_k^T)$
- After each iteration to obtain $\hat{\mathbf{L}}_{k+1}$, we can estimate Γ_{k+1} and Δ_{k+1} by using the first-order conditions.
 - More specifically, after estimating $\hat{\mathbf{L}}_{k+1}$, the objective function is as the quadratic form wrt. Γ_k, Δ_k , so we additionally minimize it and renew $\Gamma_{k+1}, \Delta_{k+1}$.
 - Finally, replace the $\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \Gamma_k \mathbf{1}_T^T - \mathbf{1}_N \Delta_k^T)$ with $\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \Gamma_{k+1} \mathbf{1}_T^T - \mathbf{1}_N \Delta_{k+1}^T)$, then proceed the algorithm to obtain $\hat{\mathbf{L}}_{k+2}$.
- we can interpret the term $\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \Gamma_k \mathbf{1}_T^T - \mathbf{1}_N \Delta_k^T)$ as an **invariant (fixed)** factor term.
- Actually, this separation of two FEs makes practical imputation greatly improved.

- The optimal value of λ is selected through **cross-validation**.
- The theory of asymptotic distribution of $\mathbf{L}^* - \hat{\mathbf{L}}$ to construct CI has not yet developed.
- Instead, using resampling method, we see the fluctuation of the imputed matrix and construct CI like permutation methods in ADH-synthetic control method.
 - Randomly select the subset \mathcal{O}_k of \mathcal{O} for $k = 1, \dots, K$.
 - Then, using the sample in \mathcal{O}_k , calculate MC-NNM estimator $\hat{\mathbf{L}}^{(k)}$.
 - Finally, we construct a pointwise CI for $\{\hat{L}_{it}^{(k)}\}$.
- However, according to Choi and Yuan(2024) which is the extended literature of Athey et al.(2021), their **debiased** estimator for \mathbf{L}^* is **asymptotic normal** (but pointwisely).

The relationship with horizontal and vertical regressions

Interpretation as generalization

- For simplicity, we assume when only one missing $Y_{NT}(0)$ exists i.e. $\mathcal{M} = (N, T)$
- Let the estimand \mathbf{Y} partition as $\begin{pmatrix} \mathbf{Y}_0 & \mathbf{y}_1 \\ \mathbf{y}_2^\top & ? \end{pmatrix}$,
where $\mathbf{Y}_0 \in \mathbb{R}^{N-1} \times \mathbb{R}^{T-1}$, $\mathbf{y}_1 \in \mathbb{R}^{N-1}$, $\mathbf{y}_2 \in \mathbb{R}^{T-1}$.
- For a given positive integer R , define an $N \times R$ matrix \mathbf{A} , an $T \times R$ matrix \mathbf{B} , a N -dim. vector γ and a R -dim. vector δ , then, the objective function w.r.t. MSE is

$$Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta) = \frac{1}{|\mathcal{O}|} \|\mathbf{P}_{\mathcal{O}}(\mathbb{Y} - \mathbf{A}\mathbf{B}^\top - \gamma\mathbf{1}_T^\top - \mathbf{1}_N\delta^\top)\|_F^2$$

- However, there is no unique solution for \mathbf{A}, \mathbf{B} unless restrict them.
→ Thus, we compare the **restriction** of each methods.

- Nuclear norm matrix completion

$$\begin{aligned} & (R_{\lambda}^{\text{mc-nnm}}, \mathbf{A}_{\lambda}^{\text{mc-nnm}}, \mathbf{B}_{\lambda}^{\text{mc-nnm}}, \gamma_{\lambda}^{\text{mc-nnm}}, \delta_{\lambda}^{\text{mc-nnm}}) \\ &= \arg \min_{R, \mathbf{A}, \mathbf{B}, \gamma, \delta} \left\{ Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta) + \frac{\lambda}{2} \|\mathbf{A}\|_F^2 + \frac{\lambda}{2} \|\mathbf{B}\|_F^2 \right\}. \end{aligned}$$

- **Fact 3.** $\|\mathbf{L}\|_* = \min_{\mathbf{A}, \mathbf{B}: \mathbf{L} = \mathbf{A}\mathbf{B}'} \frac{1}{2} (\|\mathbf{A}\|_F^2 + \|\mathbf{B}\|_F^2).$
- In second and third terms, we regularize \mathbf{A}, \mathbf{B} so that the minimization problem has a unique solution.
→ Compared with other methods, there is a big difference in that MC-NNM does not restrict the form of the matrix but just regularize it in a data-driven manner.

Horizontal regression (thin matrix, $N > T$)

- Horizontal regression estimator, defined if $N > T$: thin matrix

$$(R^{hr}, \mathbf{A}^{hr}, \mathbf{B}^{hr}, \gamma^{hr}, \delta^{hr}) = \arg \min_{R, \mathbf{A}, \mathbf{B}, \gamma, \delta} Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta),$$

$$\text{subject to } R = T - 1, \quad \mathbf{A} = \begin{pmatrix} \mathbf{Y}_0 \\ \mathbf{y}_2^\top \end{pmatrix}, \quad \gamma = 0, \quad \delta_1 = \dots = \delta_{T-1} = 0.$$

- The solution for \mathbf{B} is

$$\mathbf{B}^{hr\top} = \begin{pmatrix} E_{T-1} & \hat{\beta} \end{pmatrix}, \quad (\hat{\beta}, \hat{\delta}_T) = \arg \min_{\beta, \delta_T} \sum_{i=1}^{N-1} \left(Y_{iT} - \delta_T - \sum_{t=1}^{T-1} \beta_t Y_{it} \right)^2.$$

- Restrict the form of \mathbf{A} and (of course) assume no individual effect ($\gamma = 0$).

Vertical regression (fat matrix, $T > N$)

- Vertical regression estimator, defined if $T > N$: fat matrix

$$(R^{vt}, \mathbf{A}^{vt}, \mathbf{B}^{vt}, \gamma^{vt}, \delta^{vt}) = \arg \min_{R, \mathbf{A}, \mathbf{B}, \gamma, \delta} Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta),$$

$$\text{subject to } R = N - 1, \quad \mathbf{B}^\top = \begin{pmatrix} \mathbf{Y}_0 & \mathbf{y}_1 \end{pmatrix}, \quad \gamma_1 = \cdots = \gamma_{N-1} = 0, \quad \delta = 0.$$

- The solution for \mathbf{A} is

$$\mathbf{A}^{vt} = \begin{pmatrix} E_{N-1} \\ \hat{\alpha} \end{pmatrix}, \quad (\hat{\alpha}, \hat{\gamma}_N) = \arg \min_{\alpha, \gamma_N} \sum_{t=1}^{T-1} \left(Y_{Nt} - \gamma_N - \sum_{i=1}^{N-1} \alpha_i Y_{it} \right)^2.$$

- Restrict the form of \mathbf{B} and (of course) assume no time effect ($\delta = 0$).

(iv) (synthetic control),

$$(R^{\text{sc-adh}}, \mathbf{A}^{\text{sc-adh}}, \mathbf{B}^{\text{sc-adh}}, \gamma^{\text{sc-adh}}, \delta^{\text{sc-adh}}) \\ = \underset{R, \mathbf{A}, \mathbf{B}, \gamma, \delta}{\operatorname{argmin}} Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta),$$

subject to

$$R = N - 1, \quad \mathbf{B} = \begin{pmatrix} \mathbf{Y}_0^\top \\ \mathbf{y}_1^\top \end{pmatrix}, \quad \delta = 0, \quad \gamma = 0, \\ \forall i, A_{iT} \geq 0, \quad \sum_{i=1}^{N-1} A_{iT} = 1,$$

(v) (vertical regression, elastic net),

$$(R^{\text{vt-en}}, \mathbf{A}^{\text{vt-en}}, \mathbf{B}^{\text{vt-en}}, \gamma^{\text{vt-en}}, \delta^{\text{vt-en}}) \\ = \underset{R, \mathbf{A}, \mathbf{B}, \gamma, \delta}{\operatorname{argmin}} \left\{ Q(\mathbf{Y}; R, \mathbf{A}, \mathbf{B}, \gamma, \delta) \right. \\ \left. + \lambda \left[\frac{1-\alpha}{2} \left\| \begin{pmatrix} \mathbf{a}_2 \\ \mathbf{a}_3 \end{pmatrix} \right\|_F^2 + \alpha \left\| \begin{pmatrix} \mathbf{a}_2 \\ \mathbf{a}_3 \end{pmatrix} \right\|_1 \right] \right\},$$

subject to

$$R = N - 1, \quad \mathbf{B} = \begin{pmatrix} \mathbf{Y}_0^\top \\ \mathbf{y}_1^\top \end{pmatrix}, \\ \gamma_1 = \gamma_2 = \dots = \gamma_{N-1} = 0, \quad \delta = 0,$$

where \mathbf{A} is partitioned as

$$\mathbf{A} = \begin{pmatrix} \tilde{\mathbf{A}} & \mathbf{a}_1 \\ \mathbf{a}_2^\top & \mathbf{a}_3 \end{pmatrix},$$

Theoretical Bounds for the Estimation Error

Estimation error

- Define p_c to be a minimum expected proportion of control (never-treated) unit.

$$p_c \equiv \min_{1 \leq i \leq N} \pi_T^{(i)}, \text{ where } \pi_T^{(i)} = \mathbb{P}(t_i = T)$$

Assumption 2

- p_c is sufficiently large even as $N, T \rightarrow \infty$ with $N \geq T$. (remained control units)

$$p_c \gtrsim \frac{\log^{3/2}(N+T)}{\sqrt{T}} \vee \frac{\sqrt{T} \log^{3/2}(N+T)}{N}$$

- True matrix \mathbf{L}^* is a low rank.

- Under this assumption, $\frac{\|\mathbf{L}^* - \hat{\mathbf{L}}\|_F}{\sqrt{NT}}$ has an upper bound with a high probability $(\mathbb{P} = 1 - \frac{2}{(N+T)^2})$, then, MC-NNM estimator has a **consistency**.


Empirical application

Comparing each method

- In a real data matrix \mathbf{Y} where **no units is treated**, we choose units as "hypothetical treated" units (=regard as missing) and aim to predict (impute) their value.
→ Technically, compare the average root-mean-squared-error(RMSE) to assess which of the algorithm generally perform well.
- The compared estimators (algorithms) are as follows.
 - DiD : 2 回差分により counterfactual を補完する
 - VT-EN : elastics net で vertical regression function を作り、その推定値で補完
 - HR-EN : elastics net で horizontal regression function を作り、その推定値で補完
 - SC-ADH : classical approach to construct "synthetic control".
 - MC-NNM : 本日の主役

ADH(2010): California smoking data

- $N = 38, T = 31$
 - Simultaneous adoption : Let randomly selected $N_t = 8$ units "treated" in period $T_0 + 1$.
 - Staggered adoption : Let randomly selected $N_t = 35$ units "treated" in some periods after period T .



- (a) : VT-EN performs well on the whole, DiD is poor.
- (b) : MC-NNM is superior (large sample is favorable!), and VT-EN is generally poor.

Final marks

Summary and Conclusion

- Matrix completion is a imputing method for the missing $Y_{it}(0)$ where $(i, t) \in \mathcal{M}$.
- MC-NNM estimator generalizes and nests many estimators such as DiD, ADH-SC, vertical or horizontal (penalized) regression, and interactive fixed effect model.
- The critical difference with previous DiD estimators is that MC-NNM **holistically regularizes** latent factors through nuclear-norm minimization which induces a low-rank matrix (sparsity).
 - Intrinsic factor vs. Explicitly imposed factor
- Under appropriate conditions (low-rank structure, sufficient number of control units, sufficiently large N, T), the MC-NNM estimator achieves **consistency**.
- Practically, this estimator performs well with large N and T , and allows for a relatively large number of factors.

Further extensions

- To consider the model with covariates, we can separately include them in $\frac{1}{|\mathcal{O}|} \|\mathbf{P}_{\mathcal{O}}(\mathbf{Y} - \mathbf{A}\mathbf{B}^{\top} - \gamma\mathbf{1}_T^{\top} - \mathbf{1}_N\delta^{\top})\|_F^2$.
- For further advanced developments, a recent study by Choi and Yuan (2024, JASA) proposes a debiased estimator for \mathbf{L}^* with rigorous asymptotic inference (pointwise asymptotic normality).
- They introduce a **cross-fitting** (sample-splitting) strategy and further propose group-based ATT estimators, extending the matrix completion literature toward causal inference with stronger theoretical guarantees.

References

References

- Susan Athey, Mohsen Bayati, Nikolay Doudchenko, Guido Imbens, and Khashayar Khosravi (2021), *Matrix Completion Methods for Causal Panel Data Models*, Journal of the American Statistical Association.
- Licheng Liu, Ye Wang, and Yiqing Xu (2024), *A Practical Guide to Counterfactual Estimators for Causal Inference with Time-Series Cross-Sectional Data*, American Journal of Political Science.
- Jungjun Choi and Ming Yuan(2024), *Matrix Completion When Missing Is Not at Random and Its Applications in Causal Panel Data Models*, Journal of the American Statistical Association.
- Martin J. Wainwright (2019), *High-Dimensional Statistics*, Cambridge University Press. (mainly Chapter 10)
- 富岡亮太 (2015), スパース性に基づく機械学習, 講談社.
- 川口康平, 澤田真行 (2024), 因果推論の計量経済学, 日本評論社.